



---

# APACHE HBASE ON DOCKER

---

**BIG DATA ANALYTICS  
FINAL PROJECT**

**NAME: SYED ASAD RIZVI**

**ERP ID: 25365**

**PROGRAM: MS – DATA SCIENCE**

Dataset Used: E-commerce Behavior Data – Multi Category Store

## **Implementation Steps:**

### **HBase Container Setup:**

- 1) C:\Users\HR Computers>f:  
F:\>cd big.data
  - *For container setup, first create a folder 'big.data' in F: drive and navigate to it in the command prompt.*
- 2) F:\big.data>git clone <https://github.com/big-data-europe/docker-hbase.git>
  - *Clone the docker-hbase folder by from the mentioned link by executing the 'git clone' command.*

Cloning into 'docker-hbase'...

remote: Enumerating objects: 146, done.

remote: Counting objects: 100% (42/42), done.

remote: Compressing objects: 100% (19/19), done.

Receiving objects: 63% (92/146)used 23 (delta 23), pack-reused 104Receiving objects: 57% (84/146)

Receiving objects: 100% (146/146), 25.96 KiB | 604.00 KiB/s, done.

Resolving deltas: 100% (82/82), done.

- 3) F:\big.data>cd docker-hbase  
F:\big.data\docker-hbase>docker-compose -f docker-compose-distributed-local.yml up -d
  - *Navigate to the created folder 'docker-hbase' and run the docker-compose command to finally create the 'HBase container' in the distributed mode.*

[+] Running 12/12

- Network docker-hbase\_default Created

1.05

- Volume "docker-hbase\_hadoop\_historyserver" Created

0.15

- Volume "docker-hbase\_hadoop\_namenode" Created

0.1s

- Volume "docker-hbase\_hadoop\_datanode" Created

0.1s

- Container zoo Started

22.4s

- Container datanode Started

18.7s

- Container hbase-regionserver Started

16.2s

- Container nodemanager Started

23.4s

- Container namenode Started

12.7s

- Container resourcemanager Started

23.5s

- Container historyserver Started

17.3s

- Container hbase-master Started

14.1s

## Creating Bridge Network (Multi-node):

1) F:\big.data\docker-hbase>docker network ls

- Upon creating the HBase container, a docker bridge network named 'docker-hbase\_default' is created automatically.

NETWORK ID	NAME	DRIVER	SCOPE
4c3f1ab6301a	bridge	bridge	local
e25448a9d187	docker-hbase_default	bridge	local
b8de1f6c79f3	host	host	local
8b9a8940e314	none	null	local

2) F:\big.data\docker-hbase>docker network inspect docker-hbase\_default

- The bridge network 'docker-hbase\_default' contains all the containers running inside HBase (such as datanode, namenode etc.), which are connected like a bridge, so that they can communicate with each other.

```
[
  {
    "Name": "docker-hbase_default",
    "Id": "acd6f64810f937dd6dd72acboa7786c609abo8f2a4c64f13552c5ccfd2boe923",
    "Created": "2022-05-27T13:48:37.5925704Z",
    "Scope": "local",
    "Driver": "bridge",
    "EnableIPv6": false,
    "IPAM": {
      "Driver": "default",
      "Options": null,
      "Config": [
        {
          "Subnet": "172.18.0.0/16",
          "Gateway": "172.18.0.1"
        }
      ]
    },
    "Internal": false,
    "Attachable": false,
    "Ingress": false,
    "ConfigFrom": {
      "Network": ""
    },
    "ConfigOnly": false,
```

**"Containers": {**

```

    "2721065a329f9fa7a98a7a4274c7150a929d2bb40e671620a63375603753ceb1": {
      "Name": "nodemanager",
      "EndpointID": "316feae367fe7dac3d777dob73f5074ec151fbe1db84138a7a6c7a788cf8358c",
      "MacAddress": "02:42:ac:12:00:08",
      "IPv4Address": "172.18.0.8/16",
      "IPv6Address": ""
    },
    "77bb5c2f0f9bd5e57d5df03e8a29dab662bafddc8444d80512437721c97d20ae": {
      "Name": "historyserver",
      "EndpointID": "5d7b2ea9f4bd5e53669a5e32b46bedd7137bc331f1cb9155aa2e1cc2b4d975d1",
      "MacAddress": "02:42:ac:12:00:05",
      "IPv4Address": "172.18.0.5/16",
      "IPv6Address": ""
    },
    "9f611b702bb3ad392b9ef00ed9d9e7b2bfbfc2e68eafe171d77b833fb42115b3": {
      "Name": "hbase-regionserver",
      "EndpointID": "23788a15404c14d31a45ea5ec452e100c905b74a17ef8a28bfeaa3010f6e080b",
      "MacAddress": "02:42:ac:12:00:04",
      "IPv4Address": "172.18.0.4/16",
      "IPv6Address": ""
    },
    "b2e026235fda84eb19d7a737e3651c6063ee1ff3a861f6bdafdocea49afbedac": {
      "Name": "resourcemanager",
      "EndpointID": "45d76d887722b83afa6a0cc1d2foaf2cf1568c68f2392fc50d4b6a2ee510fad",
      "MacAddress": "02:42:ac:12:00:09",
      "IPv4Address": "172.18.0.9/16",
      "IPv6Address": ""
    },
    "d18db1728f63db4d917a48b90856eca279c1e69ef26eb15593cf1885edo6bdfa": {
      "Name": "zoo",
      "EndpointID": "16406f811a9c1086dd7e79dfccad1e565213fb555dab03f7c1d5cfd66761a0c6",
      "MacAddress": "02:42:ac:12:00:07",
      "IPv4Address": "172.18.0.7/16",
      "IPv6Address": ""
    },
    "focd77662772e61daocd7c9359147aeac271af5df7b492325a16687d4c909a4f": {
      "Name": "namenode",
      "EndpointID": "00aa495441a4e13b8f6c78doef5b8906db16fc65df576749e70b10dof799f9geb",
      "MacAddress": "02:42:ac:12:00:02",
      "IPv4Address": "172.18.0.2/16",
      "IPv6Address": ""
    },
    "f9216f7f4f5c322e9a37odddcb9e6a9b4597878dc1de6bf1889cf778b761d13e": {
      "Name": "hbase-master",
      "EndpointID": "2c5ebbb94e66afa15a4dcf46ec051bcd3c736e09d31c8ad9db89e8ea1e7d9ce",
      "MacAddress": "02:42:ac:12:00:03",
      "IPv4Address": "172.18.0.3/16",
      "IPv6Address": ""
    }
  }

```

```
    },
    "fbfd495a4588a8f44835b81e3628042a2f69718a1fc074f54f36503d53930abb": {
      "Name": "datanode",
      "EndpointID": "24602bbe2dbcab193acdoob7bbd1fde6cbaaea92802d1a11715932f03abdb7fe",
      "MacAddress": "02:42:ac:12:00:06",
      "IPv4Address": "172.18.0.6/16",
      "IPv6Address": ""
    }
  },
  "Options": {},
  "Labels": {
    "com.docker.compose.network": "default",
    "com.docker.compose.project": "docker-hbase",
    "com.docker.compose.version": "2.2.3"
  }
}
```

## Importing Dataset:

- **Importing CSV from local system to Hadoop container.**

1) F:\big.data\docker-hbase>docker cp ecommerce\_data.csv f0cd77662772:/hadoop-data

2) F:\big.data\docker-hbase>docker exec -it f0cd77662772 /bin/bash

3) root@f0cd77662772:/# cd hadoop-data

root@f0cd77662772:/hadoop-data# ls

[ecommerce\\_data.csv](#)

- **Copying CSV from container to HDFS.**

1) root@f0cd77662772:/hadoop-data# hadoop fs -mkdir -p /bda/hbasedata

2) root@f0cd77662772:/hadoop-data# hadoop fs -copyFromLocal /hadoop-data/\*.csv /bda/hbasedata

3) root@f0cd77662772:/hadoop-data# hadoop fs -ls /bda/hbasedata

[Found 1 items](#)

[-rw-r--r-- 3 root supergroup 5668612855 2022-05-27 14:14 /bda/hbasedata/ecommerce\\_data.csv](#)

- **Importing CSV from local system to HBase distributed container.**

1) F:\big.data\docker-hbase>docker cp ecommerce\_data.csv f9216f7f4f5c:/hadoop-data

2) F:\big.data\docker-hbase>docker exec -it f9216f7f4f5c /bin/bash

3) root@hbase-master:/# cd hadoop-data

root@hbase-master:/hadoop-data# ls

[ecommerce\\_data.csv](#)

- **Entering in Hbase container's command line/shell.**

4) root@hbase-master:~# hbase shell

2022-05-27 14:30:24,073 WARN [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable  
HBase Shell; enter 'help<RETURN>' for list of supported commands.  
Type "exit<RETURN>" to leave the HBase Shell  
Version 1.2.6, rUnknown, Mon May 29 02:25:32 CDT 2017



## Table Creation & Data Mapping:

1) hbase(main):006:0> hbase(main):009:0> create  
'ecommerce2','event\_time','cf1','cf2','cf3','cf4','cf5','cf6','cf7','cf8'

- Creates a table named 'ecommerce2' with **event\_time** as row key, along with eight column families (cf1 to cf8).

o row(s) in 4.6490 seconds

=> Hbase::Table - ecommerce2

2) hbase(main):003:0> create 'hbase:labels', 'f'

- Creates another table to avoid a recurring error of HBase.

o row(s) in 4.5770 seconds

=> Hbase::Table - hbase:labels

3) root@hbase-master:/# hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -  
Dimporttsv.separator=';' -  
Dimporttsv.columns=HBASE\_ROW\_KEY,cf1:event\_type,cf1:product\_id,cf2:category\_id,c  
f2:category\_code,cf3:brand,cf3:price,cf4:user\_id,cf4:user\_session ecommerce /hadoop-  
data/ecommerce\_data.csv

- After table creation and data preparation, now mapping and loading the data from HDFS to HBase.

```
.....
.....
.....
2022-05-27 17:47:26,802 INFO [LocalJobRunner Map Task Executor #0]
client.ConnectionManager$HConnectionImplementation: Closing zookeeper sessionId=0x18105c71a2e02ac
2022-05-27 17:47:26,957 INFO [LocalJobRunner Map Task Executor #0-EventThread] zookeeper.ClientCnxn:
EventThread shut down
2022-05-27 17:47:26,957 INFO [LocalJobRunner Map Task Executor #0] zookeeper.ZooKeeper: Session:
0x18105c71a2e02ac closed
2022-05-27 17:47:26,967 INFO [LocalJobRunner Map Task Executor #0] mapred.Task:
Task:attempt_local330161248_0001_m_000168_0 is done. And is in the process of committing
2022-05-27 17:47:26,968 INFO [LocalJobRunner Map Task Executor #0] mapred.LocalJobRunner: map
```

```
2022-05-27 17:47:26,968 INFO [LocalJobRunner Map Task Executor #0] mapred.Task: Task
'attempt_local330161248_0001_m_000168_o' done.
2022-05-27 17:47:26,968 INFO [LocalJobRunner Map Task Executor #0] mapred.LocalJobRunner: Finishing
task: attempt_local330161248_0001_m_000168_o
2022-05-27 17:47:26,968 INFO [Thread-33] mapred.LocalJobRunner: map task executor complete.
2022-05-27 17:47:27,256 INFO [main] mapreduce.Job: Job job_local330161248_0001 completed successfully
2022-05-27 17:47:27,413 INFO [main] mapreduce.Job: Counters: 24
  File System Counters
    FILE: Number of bytes read=486465058211
    FILE: Number of bytes written=4425411185
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=0
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=0
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=0
  Map-Reduce Framework
    Map input records=42448765
    Map output records=42448765
    Input split bytes=17069
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=24817
    CPU time spent (ms)=0
    Physical memory (bytes) snapshot=0
    Virtual memory (bytes) snapshot=0
    Total committed heap usage (bytes)=34267856896
  ImportTsv
    Bad Lines=0
  File Input Format Counters
    Bytes Read=5669300983
  File Output Format Counters
    Bytes Written=0
```

## Table Management Commands:

1) hbase(main):007:0> list

- Lists all the tables in HBase to check if the table 'ecommerce2' is present.

TABLE

ecommerce2

1 row(s) in 0.6370 seconds

=> ["ecommerce2"]

2) hbase(main):002:0> describe 'ecommerce2'

- Gives information about the column families present in the table 'ecommerce2'.

Table ecommerce2 is ENABLED

ecommerce2

COLUMN FAMILIES DESCRIPTION

{NAME => 'cf1', BLOOMFILTER => 'ROW', VERSIONS => '1', IN\_MEMORY => 'false', KEEP\_DELETED\_CELLS => 'FALSE', DATA\_BLOCK\_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN\_VERSIONS => '1'}

o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION\_SCOPE => '0'}

{NAME => 'cf2', BLOOMFILTER => 'ROW', VERSIONS => '1', IN\_MEMORY => 'false', KEEP\_DELETED\_CELLS => 'FALSE', DATA\_BLOCK\_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN\_VERSIONS => '1'}

o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION\_SCOPE => '0'}

{NAME => 'cf3', BLOOMFILTER => 'ROW', VERSIONS => '1', IN\_MEMORY => 'false', KEEP\_DELETED\_CELLS => 'FALSE', DATA\_BLOCK\_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN\_VERSIONS => '1'}

o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION\_SCOPE => '0'}

{NAME => 'cf4', BLOOMFILTER => 'ROW', VERSIONS => '1', IN\_MEMORY => 'false', KEEP\_DELETED\_CELLS => 'FALSE', DATA\_BLOCK\_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN\_VERSIONS => '1'}

o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION\_SCOPE => '0'}

{NAME => 'cf5', BLOOMFILTER => 'ROW', VERSIONS => '1', IN\_MEMORY => 'false', KEEP\_DELETED\_CELLS => 'FALSE', DATA\_BLOCK\_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN\_VERSIONS => '1'}

o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION\_SCOPE => '0'}

```
{NAME => 'cf6', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS =>
'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS
=> '
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}

{NAME => 'cf7', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS =>
'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS
=> '
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}

{NAME => 'cf8', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS =>
'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS
=> '
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}

{NAME => 'event_time', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false',
KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION =>
'NONE', MIN_VERSIONS => '1', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}

NS => 'o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}

9 row(s) in 0.1680 seconds
```

3) hbase(main):008:0> is\_enabled 'ecommerce2'

- Checks whether the table 'ecommerce2' is enabled or not.

```
true
0 row(s) in 0.0160 seconds
```

4) hbase(main):009:0> exists 'ecommerce2'

- Checks the existence of the table 'ecommerce2'.

```
Table ecommerce2 does exist
0 row(s) in 0.0130 seconds
```

5) hbase(main):011:0> is\_disabled 'ecommerce2'

- Checks whether the table 'ecommerce2' is disabled or not.

```
false
0 row(s) in 0.0630 seconds
```

```
6) hbase(main):003:0> alter 'ecommerce2', NAME => 'cf8', VERSIONS => 5
hbase(main):004:0> describe 'ecommerce2'
```

- *Changing the 'cf8' column family in table 'ecommerce2' from the current value to keep a maximum of 5 cell VERSIONS.*

Updating all regions with the new schema...

0/1 regions updated.

1/1 regions updated.

Done.

0 row(s) in 4.1910 seconds

Table ecommerce2 is ENABLED

ecommerce2

COLUMN FAMILIES DESCRIPTION

```
{NAME => 'cf1', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '1'}
```

```
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'cf2', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '1'}
```

```
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'cf3', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '1'}
```

```
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'cf4', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '1'}
```

```
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'cf5', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '1'}
```

```
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'cf6', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '1'}
```

```
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'cf7', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '1'}
```

```
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'cf8', BLOOMFILTER => 'ROW', VERSIONS => '5', IN_MEMORY => 'false', KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION => 'NONE', MIN_VERSIONS => '1'}
```

```
o', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

```
{NAME => 'event_time', BLOOMFILTER => 'ROW', VERSIONS => '1', IN_MEMORY => 'false',
KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', TTL => 'FOREVER', COMPRESSION =>
'NONE', MIN_VERSIONS => '1', BLOCKCACHE => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
9 row(s) in 0.0430 seconds
```

7) hbase(main):001:0> alter\_status 'ecommerce2'

- *Indicating the number of regions of the table 'ecommerce2' which have received the updated schema.*

1/1 regions updated.

Done.

8) hbase(main):006:0> locate\_region 'ecommerce2', '2019-10-01 00:00:30 UTC'

- *Locate the region of table 'ecommerce2' and the row key '2019-10-01 00:00:30 UTC'.*

```
HOST                                REGION
hbase-region:16020                 {ENCODED => 294ce23b90015bd3414f11c78d906d45, NAME =>
'ecommerce2,,1653667447076.294ce23b90015bd3414f11c78d906d45.', STARTKEY => "", ENDKEY => ""}
1 row(s) in 0.0240 seconds
```

9) hbase(main):006:0> disable 'ecommerce\_data'

- *Disables the table 'ecommerce\_data'. If the table needs to be dropped, it has to disable first.*

0 row(s) in 4.5210 seconds

10) hbase(main):007:0> drop 'ecommerce\_data'

- *Deletes the table 'ecommerce\_data' present in HBase.*

0 row(s) in 4.4260 seconds

## Data Manipulation Commands:

- 1) hbase(main):020:0> append 'ecommerce2', '2019-10-01 00:00:08 UTC',  
'cf4:category\_code', 'appliances.kitchen.mixer'

hbase(main):022:0> get 'ecommerce2', '2019-10-01 00:00:08 UTC', {COLUMN => 'cf4'}

- *Inserts a new value of 'appliances.kitchen.mixer' in the column family 'category\_code' against row key '2019-10-01 00:00:08 UTC', and display it.*

0 row(s) in 0.1940 seconds

COLUMN	CELL
cf4:category_code	timestamp=1654002775428, value=appliances.kitchen.mixer

1 row(s) in 0.0420 seconds

- 2) hbase(main):001:0> count 'ecommerce2', CACHE => 10000000

- *Count the number of rows in the table 'ecommerce2'. The current count is displayed every 1000 rows by default.*

Current count: 1000, row: 2019-10-01 02:19:13 UTC  
 Current count: 2000, row: 2019-10-01 02:35:54 UTC  
 Current count: 3000, row: 2019-10-01 02:52:34 UTC  
 Current count: 4000, row: 2019-10-01 03:09:14 UTC  
 Current count: 5000, row: 2019-10-01 03:25:54 UTC  
 Current count: 6000, row: 2019-10-01 03:42:34 UTC  
 Current count: 7000, row: 2019-10-01 03:59:14 UTC  
 Current count: 8000, row: 2019-10-01 04:15:54 UTC  
 Current count: 9000, row: 2019-10-01 04:32:34 UTC  
 Current count: 10000, row: 2019-10-01 04:49:14 UTC  
 Current count: 11000, row: 2019-10-01 05:05:54 UTC  
 Current count: 12000, row: 2019-10-01 05:22:34 UTC

.....  
 .....  
 .....  
 .....

Current count: 2621000, row: 2019-10-31 23:49:49 UTC  
 2621539 row(s) in 69.1020 seconds

3) hbase(main):002:0> get 'ecommerce2', '2019-10-01 00:00:22 UTC'

- *Get the columns and their values present against the row key '2019-10-01 00:00:22 UTC' in the table 'ecommerce2'.*

COLUMN	CELL
cf1:event_type	timestamp=1653667708084, value=view
cf2:product_id	timestamp=1653667708084, value=1480714
cf3:category_id	timestamp=1653667708084, value=2053013561092866779
cf4:category_code	timestamp=1653667708084, value=computers.desktop
cf5:brand	timestamp=1653667708084, value=pulser
cf6:price	timestamp=1653667708084, value=921.49
cf7:user_id	timestamp=1653667708084, value=512742880
cf8:user_session	timestamp=1653667708084, value=0d0d91c2-c9c2-4e81-90a5-86594dec0b9

8 row(s) in 0.1640 seconds

4) hbase(main):004:0> get 'ecommerce2', '2019-10-01 03:57:38 UTC', {COLUMN => 'cf6'}

- *Get the column family 'cf6' and its value present against the row key '2019-10-01 03:57:38 UTC' in the table 'ecommerce2'.*

COLUMN	CELL
cf6:price	timestamp=1653667708084, value=609.72

1 row(s) in 0.0120 seconds

5) hbase(main):006:0> get 'ecommerce2', '2019-10-01 05:50:41 UTC', {COLUMN => ['cf3','cf4','cf5']}

- *Get the column families 'cf3, cf4, cf5' and their respective values present against the row key '2019-10-01 05:50:41 UTC' in the table 'ecommerce2'.*

COLUMN	CELL
cf3:category_id	timestamp=1653667708084, value=2053013561092866779
cf4:category_code	timestamp=1653667708084, value=computers.desktop
cf5:brand	timestamp=1653667708084, value=pulser

3 row(s) in 0.0190 seconds



6) hbase(main):014:0> get 'ecommerce2', '2019-10-01 00:00:23 UTC', {FILTER => "ValueFilter(=, 'binary:midea')"} }

- Using ValueFilter, get the value 'midea' present against the row key '2019-10-01 05:50:41 UTC' in the table 'ecommerce2'.

COLUMN	CELL
cf5:brand	timestamp=1653667708084, value=midea

1 row(s) in 0.0100 seconds

7) hbase(main):010:0> get\_splits 'ecommerce2'

- Returns the split points for the table 'ecommerce2'.

Total number of splits = 1

=> []

8) hbase(main):032:0> put 'ecommerce2', '2019-10-01 00:00:15 UTC', 'cf4:category\_code', 'apparel.perfumes.men'  
hbase(main):040:0> get 'ecommerce2', '2019-10-01 00:00:15 UTC'

- Put the value 'apparel.perfumes.men' in the column 'category\_code' of family 'cf4' against the row key '2019-10-01 00:00:15 UTC', and display the contents.

0 row(s) in 0.0700 seconds

COLUMN	CELL
cf1:event_type	timestamp=1653667708084, value=view
cf2:product_id	timestamp=1653667708084, value=44600062
cf3:category_id	timestamp=1653667708084, value=2103807459595387724
cf4:category_code	timestamp=1653816014112, value=apparel.perfumes.men
cf5:brand	timestamp=1653667708084, value=shiseido
cf6:price	timestamp=1653667708084, value=35.79
cf7:user_id	timestamp=1653667708084, value=541312140
cf8:user_session	timestamp=1653667708084, value=72d76fde-8bb3-4e00-8c23-a032dfed738c

8 row(s) in 0.0420 seconds

## HBase Filters:

9) hbase(main):001:0> show\_filters

- *Returns different filters that can be used to view the data.*

DependentColumnFilter  
KeyOnlyFilter  
ColumnCountGetFilter  
SingleColumnValueFilter  
PrefixFilter  
SingleColumnValueExcludeFilter  
FirstKeyOnlyFilter  
ColumnRangeFilter  
TimestampsFilter  
FamilyFilter  
QualifierFilter  
ColumnPrefixFilter  
RowFilter  
MultipleColumnPrefixFilter  
InclusiveStopFilter  
PageFilter  
ValueFilter  
ColumnPaginationFilter

10) hbase(main):002:0> scan 'ecommerce2', {FILTER => "KeyOnlyFilter()"}

- *Returns the key portion i.e. column name of each key-value pair in the table 'ecommerce2'.*

### ROW

2019-10-01 00:00:00 UTC  
2019-10-01 00:00:00 UTC  
2019-10-01 00:00:00 UTC  
2019-10-01 00:00:00 UTC  
2019-10-01 00:00:00 UTC  
2019-10-01 00:00:00 UTC  
2019-10-01 00:00:00 UTC  
2019-10-01 00:00:00 UTC  
2019-10-01 00:00:01 UTC  
2019-10-01 00:00:01 UTC  
2019-10-01 00:00:01 UTC  
2019-10-01 00:00:01 UTC  
2019-10-01 00:00:01 UTC  
2019-10-01 00:00:01 UTC  
2019-10-01 00:00:01 UTC

### COLUMN+CELL

column=cf1:event\_type, timestamp=1653667708084, value=  
column=cf2:product\_id, timestamp=1653667708084, value=  
column=cf3:category\_id, timestamp=1653667708084, value=  
column=cf4:category\_code, timestamp=1653667708084, value=  
column=cf5:brand, timestamp=1653667708084, value=  
column=cf6:price, timestamp=1653667708084, value=  
column=cf7:user\_id, timestamp=1653667708084, value=  
column=cf8:user\_session, timestamp=1653667708084, value=  
column=cf1:event\_type, timestamp=1653667708084, value=  
column=cf2:product\_id, timestamp=1653667708084, value=  
column=cf3:category\_id, timestamp=1653667708084, value=  
column=cf4:category\_code, timestamp=1653667708084, value=  
column=cf5:brand, timestamp=1653667708084, value=  
column=cf6:price, timestamp=1653667708084, value=

```
2019-10-01 00:00:01 UTC      column=cf7:user_id, timestamp=1653667708084, value=
2019-10-01 00:00:01 UTC      column=cf8:user_session, timestamp=1653667708084, value=
```

```
.....
.....
.....
```

11) hbase(main):001:0> scan 'ecommerce2', {FILTER => "FirstKeyOnlyFilter()"}

- Returns the first key-value pair i.e. the first column and its value against each row key in the table 'ecommerce2'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:01 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:04 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:05 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:08 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:10 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:11 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:13 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:15 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:16 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:17 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:18 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:19 UTC	column=cf1:event_type, timestamp=1653667708084, value=view

```
.....
.....
.....
```

12) hbase(main):001:0> scan 'ecommerce2', {FILTER => "PrefixFilter('1')"}

- Returns only those key-values present in a row that starts with the row prefix '1'. Since there is no such row, hence it will return empty.

ROW	COLUMN+CELL
0 row(s) in 0.4750 seconds	

13) hbase(main):002:0> scan 'ecommerce2', {FILTER => "PrefixFilter('2')"}

- Returns only those key-values present in a row that starts with the row prefix '2'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:00 UTC	column=cf2:product_id, timestamp=1653667708084, value=3900821
2019-10-01 00:00:00 UTC value=2053013552326770905	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:00 UTC value=appliances.environment.water_heater	column=cf4:category_code, timestamp=1653667708084,
2019-10-01 00:00:00 UTC	column=cf5:brand, timestamp=1653667708084, value=aqua
2019-10-01 00:00:00 UTC	column=cf6:price, timestamp=1653667708084, value=33.20
2019-10-01 00:00:00 UTC	column=cf7:user_id, timestamp=1653667708084, value=554748717
2019-10-01 00:00:00 UTC b87a-4708-9857-6336556bofcc	column=cf8:user_session, timestamp=1653667708084, value=9333dfbd-
2019-10-01 00:00:01 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:01 UTC	column=cf2:product_id, timestamp=1653667708084, value=1307067
2019-10-01 00:00:01 UTC value=2053013558920217191	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:01 UTC value=computers.notebook	column=cf4:category_code, timestamp=1653667708084,
2019-10-01 00:00:01 UTC	column=cf5:brand, timestamp=1653667708084, value=lenovo
2019-10-01 00:00:01 UTC	column=cf6:price, timestamp=1653667708084, value=251.74
2019-10-01 00:00:01 UTC	column=cf7:user_id, timestamp=1653667708084, value=550050854
2019-10-01 00:00:01 UTC 0e80-4590-96f3-13c02c18c713	column=cf8:user_session, timestamp=1653667708084, value=7c90fc70-
.....	
.....	
.....	

14) hbase(main):001:0> scan 'ecommerce2', {FILTER => "ColumnPrefixFilter('u')"}

- Returns only those key-values of a column that starts with the prefix 'u', which are 'user\_id' and 'user\_session'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf7:user_id, timestamp=1653667708084, value=554748717
2019-10-01 00:00:00 UTC b87a-4708-9857-6336556bofcc	column=cf8:user_session, timestamp=1653667708084, value=9333dfbd-
2019-10-01 00:00:01 UTC	column=cf7:user_id, timestamp=1653667708084, value=550050854
2019-10-01 00:00:01 UTC 0e80-4590-96f3-13c02c18c713	column=cf8:user_session, timestamp=1653667708084, value=7c90fc70-
2019-10-01 00:00:04 UTC	column=cf7:user_id, timestamp=1653667708084, value=535871217
2019-10-01 00:00:04 UTC 2748-4c56-95b4-8cec9ff8b80d	column=cf8:user_session, timestamp=1653667708084, value=c6bd7419-
2019-10-01 00:00:05 UTC	column=cf7:user_id, timestamp=1653667708084, value=512742880
2019-10-01 00:00:05 UTC c9c2-4e81-90a5-86594decodb9	column=cf8:user_session, timestamp=1653667708084, value=0d0d91c2-
2019-10-01 00:00:08 UTC	column=cf7:user_id, timestamp=1653667708084, value=550978835

2019-10-01 00:00:08 UTC	column=cf8:user_session, timestamp=1653667708084, value=6280d577-25c8-4147-99a7-abc6048498d6
2019-10-01 00:00:10 UTC	column=cf7:user_id, timestamp=1653667708084, value=520571932
2019-10-01 00:00:10 UTC	column=cf8:user_session, timestamp=1653667708084, value=ac1cd4e5-a3ce-4224-a2d7-ff660a105880
2019-10-01 00:00:11 UTC	column=cf7:user_id, timestamp=1653667708084, value=530282093
2019-10-01 00:00:11 UTC	column=cf8:user_session, timestamp=1653667708084, value=50a293fb-5940-41b2-baf3-17afoe812101
2019-10-01 00:00:13 UTC	column=cf7:user_id, timestamp=1653667708084, value=555444559
2019-10-01 00:00:13 UTC	column=cf8:user_session, timestamp=1653667708084, value=98b88fao-d8fa-4b9d-8a71-3dd403afab85
.....	
.....	
.....	

15) hbase(main):001:0> scan 'ecommerce2', {FILTER => "MultipleColumnPrefixFilter('b','p')"}

- Returns only those key-values present in a column that starts with any of the column prefixes 'b' and 'p', which are 'brand', 'price', and 'product\_id'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf2:product_id, timestamp=1653667708084, value=3900821
2019-10-01 00:00:00 UTC	column=cf5:brand, timestamp=1653667708084, value=aqua
2019-10-01 00:00:00 UTC	column=cf6:price, timestamp=1653667708084, value=33.20
2019-10-01 00:00:01 UTC	column=cf2:product_id, timestamp=1653667708084, value=1307067
2019-10-01 00:00:01 UTC	column=cf5:brand, timestamp=1653667708084, value=lenovo
2019-10-01 00:00:01 UTC	column=cf6:price, timestamp=1653667708084, value=251.74
2019-10-01 00:00:04 UTC	column=cf2:product_id, timestamp=1653667708084, value=1004237
2019-10-01 00:00:04 UTC	column=cf5:brand, timestamp=1653667708084, value=apple
2019-10-01 00:00:04 UTC	column=cf6:price, timestamp=1653667708084, value=1081.98
2019-10-01 00:00:05 UTC	column=cf2:product_id, timestamp=1653667708084, value=1480613
2019-10-01 00:00:05 UTC	column=cf5:brand, timestamp=1653667708084, value=pulser
2019-10-01 00:00:05 UTC	column=cf6:price, timestamp=1653667708084, value=908.62
2019-10-01 00:00:08 UTC	column=cf2:product_id, timestamp=1653667708084, value=31500053
2019-10-01 00:00:08 UTC	column=cf5:brand, timestamp=1653667708084, value=luminarc
2019-10-01 00:00:08 UTC	column=cf6:price, timestamp=1653667708084, value=41.16
.....	
.....	
.....	

16) hbase(main):001:0> scan 'ecommerce2', {FILTER => "ColumnCountGetFilter(4)"}

- Returns the first 4 rows and their columns of the table 'ecommerce2'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:00 UTC	column=cf2:product_id, timestamp=1653667708084, value=3900821
2019-10-01 00:00:00 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013552326770905	
2019-10-01 00:00:00 UTC	column=cf4:category_code, timestamp=1653667708084,
value=appliances.environment.water_heater	
2019-10-01 00:00:01 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:01 UTC	column=cf2:product_id, timestamp=1653667708084, value=1307067
2019-10-01 00:00:01 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013558920217191	
2019-10-01 00:00:01 UTC	column=cf4:category_code, timestamp=1653667708084,
value=computers.notebook	
2019-10-01 00:00:04 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:04 UTC	column=cf2:product_id, timestamp=1653667708084, value=1004237
2019-10-01 00:00:04 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013555631882655	
2019-10-01 00:00:04 UTC	column=cf4:category_code, timestamp=1653667708084,
value=electronics.smartphone	

.....

.....

.....

17) hbase(main):001:0> scan 'ecommerce2', {FILTER => "ColumnPaginationFilter(2,4)"}

- Takes two arguments, limit and offset. Returns the two columns which are present after the first four columns.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf5:brand, timestamp=1653667708084, value=aqua
2019-10-01 00:00:00 UTC	column=cf6:price, timestamp=1653667708084, value=33.20
2019-10-01 00:00:01 UTC	column=cf5:brand, timestamp=1653667708084, value=lenovo
2019-10-01 00:00:01 UTC	column=cf6:price, timestamp=1653667708084, value=251.74
2019-10-01 00:00:04 UTC	column=cf5:brand, timestamp=1653667708084, value=apple
2019-10-01 00:00:04 UTC	column=cf6:price, timestamp=1653667708084, value=1081.98
2019-10-01 00:00:05 UTC	column=cf5:brand, timestamp=1653667708084, value=pulser
2019-10-01 00:00:05 UTC	column=cf6:price, timestamp=1653667708084, value=908.62
2019-10-01 00:00:08 UTC	column=cf5:brand, timestamp=1653667708084,
value=luminarc	
2019-10-01 00:00:08 UTC	column=cf6:price, timestamp=1653667708084, value=41.16
2019-10-01 00:00:10 UTC	column=cf5:brand, timestamp=1653667708084, value=baden
2019-10-01 00:00:10 UTC	column=cf6:price, timestamp=1653667708084, value=102.71
2019-10-01 00:00:11 UTC	column=cf5:brand, timestamp=1653667708084,
value=samsung	
2019-10-01 00:00:11 UTC	column=cf6:price, timestamp=1653667708084, value=900.64

```

2019-10-01 00:00:13 UTC      column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 00:00:13 UTC      column=cf6:price, timestamp=1653667708084, value=102.38
.....
.....
.....

```

```

18) hbase(main):003:0> scan 'ecommerce2', {FILTER => "InclusiveStopFilter('2019-10-01
00:00:04 UTC')"}

```

- Returns all key-values present in rows up to the row key '2019-10-01 00:00:04 UTC', including the specified row key itself.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:00 UTC	column=cf2:product_id, timestamp=1653667708084, value=3900821
2019-10-01 00:00:00 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013552326770905	
2019-10-01 00:00:00 UTC	column=cf4:category_code, timestamp=1653667708084,
value=appliances.environment.water_heater	
2019-10-01 00:00:00 UTC	column=cf5:brand, timestamp=1653667708084, value=aqua
2019-10-01 00:00:00 UTC	column=cf6:price, timestamp=1653667708084, value=33.20
2019-10-01 00:00:00 UTC	column=cf7:user_id, timestamp=1653667708084, value=554748717
2019-10-01 00:00:00 UTC	column=cf8:user_session, timestamp=1653667708084, value=9333dfbd-
b87a-4708-9857-6336556b0fcc	
2019-10-01 00:00:01 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:01 UTC	column=cf2:product_id, timestamp=1653667708084, value=1307067
2019-10-01 00:00:01 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013558920217191	
2019-10-01 00:00:01 UTC	column=cf4:category_code, timestamp=1653667708084,
value=computers.notebook	
2019-10-01 00:00:01 UTC	column=cf5:brand, timestamp=1653667708084, value=lenovo
2019-10-01 00:00:01 UTC	column=cf6:price, timestamp=1653667708084, value=251.74
2019-10-01 00:00:01 UTC	column=cf7:user_id, timestamp=1653667708084, value=550050854
2019-10-01 00:00:01 UTC	column=cf8:user_session, timestamp=1653667708084, value=7c90fc70-
0e80-4590-96f3-13c02c18c713	
2019-10-01 00:00:04 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:04 UTC	column=cf2:product_id, timestamp=1653667708084, value=1004237
2019-10-01 00:00:04 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013555631882655	
2019-10-01 00:00:04 UTC	column=cf4:category_code, timestamp=1653667708084,
value=electronics.smartphone	
2019-10-01 00:00:04 UTC	column=cf5:brand, timestamp=1653667708084, value=apple
2019-10-01 00:00:04 UTC	column=cf6:price, timestamp=1653667708084, value=1081.98
2019-10-01 00:00:04 UTC	column=cf7:user_id, timestamp=1653667708084, value=535871217
2019-10-01 00:00:04 UTC	column=cf8:user_session, timestamp=1653667708084, value=c6bd7419-
2748-4c56-95b4-8cec9ff8b80d	







```

2019-10-01 00:00:00 UTC      column=cf6:price, timestamp=1653667708084, value=33.20
2019-10-01 00:00:00 UTC      column=cf7:user_id, timestamp=1653667708084, value=554748717
2019-10-01 00:00:00 UTC      column=cf8:user_session, timestamp=1653667708084, value=9333dfbd-
b87a-4708-9857-6336556b0fcc
2019-10-01 00:00:01 UTC      column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:01 UTC      column=cf2:product_id, timestamp=1653667708084, value=1307067
2019-10-01 00:00:01 UTC      column=cf3:category_id, timestamp=1653667708084,
value=2053013558920217191
2019-10-01 00:00:01 UTC      column=cf4:category_code, timestamp=1653667708084,
value=computers.notebook
2019-10-01 00:00:01 UTC      column=cf5:brand, timestamp=1653667708084, value=lenovo
2019-10-01 00:00:01 UTC      column=cf6:price, timestamp=1653667708084, value=251.74
2019-10-01 00:00:01 UTC      column=cf7:user_id, timestamp=1653667708084, value=550050854
2019-10-01 00:00:01 UTC      column=cf8:user_session, timestamp=1653667708084, value=7c90fc70-
0e80-4590-96f3-13c02c18c713
2019-10-01 00:00:04 UTC      column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:04 UTC      column=cf2:product_id, timestamp=1653667708084, value=1004237
2019-10-01 00:00:04 UTC      column=cf3:category_id, timestamp=1653667708084,
value=2053013555631882655
2019-10-01 00:00:04 UTC      column=cf4:category_code, timestamp=1653667708084,
value=electronics.smartphone
2019-10-01 00:00:04 UTC      column=cf5:brand, timestamp=1653667708084, value=apple
2019-10-01 00:00:04 UTC      column=cf6:price, timestamp=1653667708084, value=1081.98
2019-10-01 00:00:04 UTC      column=cf7:user_id, timestamp=1653667708084, value=535871217
2019-10-01 00:00:04 UTC      column=cf8:user_session, timestamp=1653667708084, value=c6bd7419-
2748-4c56-95b4-8cec9ff8b80d
2019-10-01 00:00:05 UTC      column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:05 UTC      column=cf2:product_id, timestamp=1653667708084, value=1480613
2019-10-01 00:00:05 UTC      column=cf3:category_id, timestamp=1653667708084,
value=2053013561092866779
2019-10-01 00:00:05 UTC      column=cf4:category_code, timestamp=1653667708084,
value=computers.desktop
2019-10-01 00:00:05 UTC      column=cf5:brand, timestamp=1653667708084, value=pulser
2019-10-01 00:00:05 UTC      column=cf6:price, timestamp=1653667708084, value=908.62
2019-10-01 00:00:05 UTC      column=cf7:user_id, timestamp=1653667708084, value=512742880
2019-10-01 00:00:05 UTC      column=cf8:user_session, timestamp=1653667708084, value=0d0d91c2-
c9c2-4e81-90a5-86594dec0db9
4 row(s) in 38.1680 seconds

```

```
21) hbase(main):006:0> scan 'ecommerce2', {FILTER => "FamilyFilter(<=, 'binaryprefix:cf3')"}

```

- Returns all the key-value pairs of column family 'cf3' and its predecessors, such that it contains column families cf1, cf2, and cf3 itself.

ROW

COLUMN+CELL

```

2019-10-01 00:00:00 UTC      column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:00 UTC      column=cf2:product_id, timestamp=1653667708084, value=3900821

```

2019-10-01 00:00:00 UTC value=2053013552326770905	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:01 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:01 UTC	column=cf2:product_id, timestamp=1653667708084, value=1307067
2019-10-01 00:00:01 UTC value=2053013558920217191	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:04 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:04 UTC	column=cf2:product_id, timestamp=1653667708084, value=1004237
2019-10-01 00:00:04 UTC value=2053013555631882655	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:05 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:05 UTC	column=cf2:product_id, timestamp=1653667708084, value=1480613
2019-10-01 00:00:05 UTC value=2053013561092866779	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:08 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:08 UTC	column=cf2:product_id, timestamp=1653667708084, value=31500053
2019-10-01 00:00:08 UTC value=2053013558031024687	column=cf3:category_id, timestamp=1653667708084,

.....

.....

.....

22) hbase(main):001:0> scan 'ecommerce2', {FILTER => "QualifierFilter(=, 'substring:and')"}{

- Returns all the key-value pairs of the columns that matches its name with the substring 'and'. So, the output column is 'brand'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf5:brand, timestamp=1653667708084, value=aqua
2019-10-01 00:00:01 UTC	column=cf5:brand, timestamp=1653667708084, value=lenovo
2019-10-01 00:00:04 UTC	column=cf5:brand, timestamp=1653667708084, value=apple
2019-10-01 00:00:05 UTC	column=cf5:brand, timestamp=1653667708084, value=pulser
2019-10-01 00:00:08 UTC value=luminarc	column=cf5:brand, timestamp=1653667708084,
2019-10-01 00:00:10 UTC	column=cf5:brand, timestamp=1653667708084, value=baden
2019-10-01 00:00:11 UTC value=samsung	column=cf5:brand, timestamp=1653667708084,
2019-10-01 00:00:13 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 00:00:15 UTC	column=cf5:brand, timestamp=1653667708084, value=shiseido
2019-10-01 00:00:16 UTC	column=cf5:brand, timestamp=1653667708084, value=brw
2019-10-01 00:00:17 UTC	column=cf5:brand, timestamp=1653667708084, value=
2019-10-01 00:00:18 UTC	column=cf5:brand, timestamp=1653667708084, value=bosch
2019-10-01 00:00:19 UTC	column=cf5:brand, timestamp=1653667708084, value=apple
2019-10-01 00:00:20 UTC	column=cf5:brand, timestamp=1653667708084, value=jbl
2019-10-01 00:00:22 UTC	column=cf5:brand, timestamp=1653667708084, value=pulser

23) hbase(main):001:0> scan 'ecommerce2', {FILTER => "QualifierFilter(=, 'substring:id')"}

- Returns all the key-value pairs of the columns that matches its name with the substring 'id'. So, the output columns are 'product\_id', 'category\_id', and 'user\_id'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf2:product_id, timestamp=1653667708084, value=3900821
2019-10-01 00:00:00 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013552326770905	
2019-10-01 00:00:00 UTC	column=cf7:user_id, timestamp=1653667708084, value=554748717
2019-10-01 00:00:01 UTC	column=cf2:product_id, timestamp=1653667708084, value=1307067
2019-10-01 00:00:01 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013558920217191	
2019-10-01 00:00:01 UTC	column=cf7:user_id, timestamp=1653667708084, value=550050854
2019-10-01 00:00:04 UTC	column=cf2:product_id, timestamp=1653667708084, value=1004237
2019-10-01 00:00:04 UTC	column=cf3:category_id, timestamp=1653667708084,
value=205301355631882655	
2019-10-01 00:00:04 UTC	column=cf7:user_id, timestamp=1653667708084, value=535871217
2019-10-01 00:00:05 UTC	column=cf2:product_id, timestamp=1653667708084, value=1480613
2019-10-01 00:00:05 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013561092866779	
2019-10-01 00:00:05 UTC	column=cf7:user_id, timestamp=1653667708084, value=512742880
2019-10-01 00:00:08 UTC	column=cf2:product_id, timestamp=1653667708084, value=31500053
2019-10-01 00:00:08 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013558031024687	
2019-10-01 00:00:08 UTC	column=cf7:user_id, timestamp=1653667708084, value=550978835

24) hbase(main):001:0> scan 'ecommerce2', {FILTER => "ValueFilter(=, 'binary:haier')"}

- Returns key-value pairs against each row that matches with the string 'haier'.

ROW	COLUMN+CELL
2019-10-01 00:00:13 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 00:13:47 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 02:41:31 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 02:51:08 UTC	column=cf5:brand, timestamp=1653667708084, value=haier

2019-10-01 02:55:39 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 02:57:46 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 02:58:39 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 02:58:52 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 03:03:11 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 03:08:10 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 03:43:14 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 03:43:15 UTC	column=cf5:brand, timestamp=1653667708084, value=haier
2019-10-01 03:53:48 UTC	column=cf5:brand, timestamp=1653667708084, value=haier

.....

.....

.....

25) hbase(main):001:0> scan 'ecommerce2', {FILTER => "DependentColumnFilter('cf4', 'category\_code')"}  
 'category\_code'")}

- Returns all the key-value pairs that have the same timestamp as in the column family 'cf4' with column 'category\_code'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:00 UTC	column=cf2:product_id, timestamp=1653667708084, value=3900821
2019-10-01 00:00:00 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013552326770905	
2019-10-01 00:00:00 UTC	column=cf4:category_code, timestamp=1653667708084,
value=appliances.environment.water_heater	
2019-10-01 00:00:00 UTC	column=cf5:brand, timestamp=1653667708084, value=aqua
2019-10-01 00:00:00 UTC	column=cf6:price, timestamp=1653667708084, value=33.20
2019-10-01 00:00:00 UTC	column=cf7:user_id, timestamp=1653667708084, value=554748717
2019-10-01 00:00:00 UTC	column=cf8:user_session, timestamp=1653667708084, value=9333dfbd-
b87a-4708-9857-6336556b0fcc	
2019-10-01 00:00:01 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:01 UTC	column=cf2:product_id, timestamp=1653667708084, value=1307067
2019-10-01 00:00:01 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013558920217191	
2019-10-01 00:00:01 UTC	column=cf4:category_code, timestamp=1653667708084,
value=computers.notebook	
2019-10-01 00:00:01 UTC	column=cf5:brand, timestamp=1653667708084, value=lenovo
2019-10-01 00:00:01 UTC	column=cf6:price, timestamp=1653667708084, value=251.74
2019-10-01 00:00:01 UTC	column=cf7:user_id, timestamp=1653667708084, value=550050854
2019-10-01 00:00:01 UTC	column=cf8:user_session, timestamp=1653667708084, value=7c90fc70-
0e80-4590-96f3-13c02c18c713	

.....

.....

.....

```
26) hbase(main):001:0> scan 'ecommerce2', {FILTER => "DependentColumnFilter('cf1',
'event_type', true)"}
```

- Returns all the key-value pairs that have the same timestamp as in the column family 'cf1' with column 'event\_type'. Family 'cf1' itself will not return because the Boolean argument 'dropDependentColumn' is set to 'true'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf2:product_id, timestamp=1653667708084, value=3900821
2019-10-01 00:00:00 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013552326770905	
2019-10-01 00:00:00 UTC	column=cf4:category_code, timestamp=1653667708084,
value=appliances.environment.water_heater	
2019-10-01 00:00:00 UTC	column=cf5:brand, timestamp=1653667708084, value=aqua
2019-10-01 00:00:00 UTC	column=cf6:price, timestamp=1653667708084, value=33.20
2019-10-01 00:00:00 UTC	column=cf7:user_id, timestamp=1653667708084, value=554748717
2019-10-01 00:00:00 UTC	column=cf8:user_session, timestamp=1653667708084, value=9333dfbd-
b87a-4708-9857-6336556bofcc	
2019-10-01 00:00:01 UTC	column=cf2:product_id, timestamp=1653667708084, value=1307067
2019-10-01 00:00:01 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013558920217191	
2019-10-01 00:00:01 UTC	column=cf4:category_code, timestamp=1653667708084,
value=computers.notebook	
2019-10-01 00:00:01 UTC	column=cf5:brand, timestamp=1653667708084, value=lenovo
2019-10-01 00:00:01 UTC	column=cf6:price, timestamp=1653667708084, value=251.74
2019-10-01 00:00:01 UTC	column=cf7:user_id, timestamp=1653667708084, value=550050854
2019-10-01 00:00:01 UTC	column=cf8:user_session, timestamp=1653667708084, value=7c90fc70-
0e80-4590-96f3-13c02c18c713	
.....	
.....	
.....	

```
27) hbase(main):002:0> scan 'ecommerce2', {FILTER => "ColumnRangeFilter('brand', true,
'product_id', true)"}
```

- Returns only those keys and columns that are between 'brand' and 'product\_id'. The Boolean arguments 'true' shows that both the specified columns are inclusive.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:00 UTC	column=cf2:product_id, timestamp=1653667708084, value=3900821
2019-10-01 00:00:00 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013552326770905	

2019-10-01 00:00:00 UTC value=appliances.environment.water_heater	column=cf4:category_code, timestamp=1653667708084,
2019-10-01 00:00:00 UTC	column=cf5:brand, timestamp=1653667708084, value=aqua
2019-10-01 00:00:00 UTC	column=cf6:price, timestamp=1653667708084, value=33.20
2019-10-01 00:00:01 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:01 UTC	column=cf2:product_id, timestamp=1653667708084, value=1307067
2019-10-01 00:00:01 UTC value=2053013558920217191	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:01 UTC value=computers.notebook	column=cf4:category_code, timestamp=1653667708084,
2019-10-01 00:00:01 UTC	column=cf5:brand, timestamp=1653667708084, value=lenovo
2019-10-01 00:00:01 UTC	column=cf6:price, timestamp=1653667708084, value=251.74
2019-10-01 00:00:04 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:04 UTC	column=cf2:product_id, timestamp=1653667708084, value=1004237
2019-10-01 00:00:04 UTC value=2053013555631882655	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:04 UTC value=electronics.smartphone	column=cf4:category_code, timestamp=1653667708084,
2019-10-01 00:00:04 UTC	column=cf5:brand, timestamp=1653667708084, value=apple
2019-10-01 00:00:04 UTC	column=cf6:price, timestamp=1653667708084, value=1081.98

.....

.....

.....

28) hbase(main):003:0> scan 'ecommerce2', {FILTER => "PrefixFilter('1') AND KeyOnlyFilter()}"}

- Returns empty, since there is no row key that starts with '1'.

ROW COLUMN+CELL  
0 row(s) in 0.0530 seconds

29) hbase(main):004:0> scan 'ecommerce2', {FILTER => "PrefixFilter('2') AND KeyOnlyFilter()}"}

- Returns only keys of those rows that starts with '2'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf1:event_type, timestamp=1653667708084, value=
2019-10-01 00:00:00 UTC	column=cf2:product_id, timestamp=1653667708084, value=
2019-10-01 00:00:00 UTC	column=cf3:category_id, timestamp=1653667708084, value=
2019-10-01 00:00:00 UTC	column=cf4:category_code, timestamp=1653667708084,
value=	



2019-10-01 00:00:00 UTC	column=cf5:brand, timestamp=1653667708084, value=
2019-10-01 00:00:00 UTC	column=cf6:price, timestamp=1653667708084, value=
2019-10-01 00:00:00 UTC	column=cf7:user_id, timestamp=1653667708084, value=
2019-10-01 00:00:00 UTC	column=cf8:user_session, timestamp=1653667708084, value=
2019-10-01 00:00:01 UTC	column=cf1:event_type, timestamp=1653667708084, value=
2019-10-01 00:00:01 UTC	column=cf2:product_id, timestamp=1653667708084, value=
2019-10-01 00:00:01 UTC	column=cf3:category_id, timestamp=1653667708084, value=
2019-10-01 00:00:01 UTC	column=cf4:category_code, timestamp=1653667708084, value=
2019-10-01 00:00:01 UTC	column=cf5:brand, timestamp=1653667708084, value=
2019-10-01 00:00:01 UTC	column=cf6:price, timestamp=1653667708084, value=
2019-10-01 00:00:01 UTC	column=cf7:user_id, timestamp=1653667708084, value=
2019-10-01 00:00:01 UTC	column=cf8:user_session, timestamp=1653667708084, value=
.....	
.....	
.....	

30) hbase(main):005:0> scan 'ecommerce2', {FILTER => "RowFilter(<, 'binary:2019-10-01 00:00:13 UTC') AND QualifierFilter(=, 'substring:user')"} }

- Returns key-value pairs of columns that starts with the string 'user' against all rows that are present before '2019-10-01 00:00:13 UTC'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf7:user_id, timestamp=1653667708084, value=554748717
2019-10-01 00:00:00 UTC	column=cf8:user_session, timestamp=1653667708084, value=9333dfbd-
b87a-4708-9857-6336556b0fcc	
2019-10-01 00:00:01 UTC	column=cf7:user_id, timestamp=1653667708084, value=550050854
2019-10-01 00:00:01 UTC	column=cf8:user_session, timestamp=1653667708084, value=7c90fc70-
0e80-4590-96f3-13c02c18c713	
2019-10-01 00:00:04 UTC	column=cf7:user_id, timestamp=1653667708084, value=535871217
2019-10-01 00:00:04 UTC	column=cf8:user_session, timestamp=1653667708084, value=c6bd7419-
2748-4c56-95b4-8cec9ff8b80d	
2019-10-01 00:00:05 UTC	column=cf7:user_id, timestamp=1653667708084, value=512742880
2019-10-01 00:00:05 UTC	column=cf8:user_session, timestamp=1653667708084, value=0d0d91c2-
c9c2-4e81-90a5-86594dec0db9	
2019-10-01 00:00:08 UTC	column=cf7:user_id, timestamp=1653667708084, value=550978835
2019-10-01 00:00:08 UTC	column=cf8:user_session, timestamp=1653667708084, value=6280d577-
25c8-4147-99a7-abc6048498d6	
2019-10-01 00:00:10 UTC	column=cf7:user_id, timestamp=1653667708084, value=520571932
2019-10-01 00:00:10 UTC	column=cf8:user_session, timestamp=1653667708084, value=ac1cd4e5-
a3ce-4224-a2d7-ff660a105880	
2019-10-01 00:00:11 UTC	column=cf7:user_id, timestamp=1653667708084, value=530282093
2019-10-01 00:00:11 UTC	column=cf8:user_session, timestamp=1653667708084, value=50a293fb-
5940-41b2-baf3-17afoe812101	

31) hbase(main):001:0> scan 'ecommerce2', {FILTER => "ValueFilter(=, 'binary:purchase') OR ValueFilter(=, 'binary:ariston')"} }

- Returns key-value pairs against each row that has either the value 'purchase' or 'ariston' in it.

ROW	COLUMN+CELL
2019-10-01 00:01:25 UTC	column=cf5:brand, timestamp=1653667708084, value=ariston
2019-10-01 00:01:28 UTC	column=cf5:brand, timestamp=1653667708084, value=ariston
2019-10-01 00:02:14 UTC	column=cf1:event_type, timestamp=1653667708084, value=purchase
2019-10-01 00:09:26 UTC	column=cf1:event_type, timestamp=1653667708084, value=purchase
2019-10-01 00:10:56 UTC	column=cf1:event_type, timestamp=1653667708084, value=purchase
2019-10-01 00:12:14 UTC	column=cf1:event_type, timestamp=1653667708084, value=purchase
2019-10-01 00:14:14 UTC	column=cf1:event_type, timestamp=1653667708084, value=purchase
2019-10-01 00:14:20 UTC	column=cf5:brand, timestamp=1653667708084, value=ariston
2019-10-01 00:16:40 UTC	column=cf5:brand, timestamp=1653667708084, value=ariston
2019-10-01 00:17:53 UTC	column=cf5:brand, timestamp=1653667708084, value=ariston
2019-10-01 02:19:12 UTC	column=cf1:event_type, timestamp=1653667708084, value=purchase
2019-10-01 02:19:34 UTC	column=cf1:event_type, timestamp=1653667708084, value=purchase
.....	.....
.....	.....
.....	.....

32) hbase(main):001:0> scan 'ecommerce2', {FILTER => "PageFilter(3) AND ColumnPrefixFilter('e')"} }

- Returns the first three rows of the table containing only those key-values of a column that starts with 'e'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:01 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:04 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
3 row(s) in 0.4050 seconds	

33) hbase(main):001:0> scan 'ecommerce2', {COLUMNS => 'cf5', FILTER => "ValueFilter(=, 'regexstring:^(d.\*)")"} }



- Returns key values of the column family 'cf5:brand' in which the name starts with the letter 'd'.

ROW	COLUMN+CELL
2019-10-01 00:00:30 UTC	column=cf5:brand, timestamp=1653667708084, value=dauscher
2019-10-01 00:04:01 UTC	column=cf5:brand, timestamp=1653667708084, value=dewalt
2019-10-01 00:04:53 UTC	column=cf5:brand, timestamp=1653667708084, value=dior
2019-10-01 00:09:04 UTC	column=cf5:brand, timestamp=1653667708084, value=delonghi
2019-10-01 00:13:35 UTC	column=cf5:brand, timestamp=1653667708084, value=dauscher
2019-10-01 00:18:58 UTC	column=cf5:brand, timestamp=1653667708084, value=doogee
2019-10-01 00:19:01 UTC	column=cf5:brand, timestamp=1653667708084, value=doogee
2019-10-01 00:19:05 UTC	column=cf5:brand, timestamp=1653667708084, value=doogee
2019-10-01 00:58:07 UTC	column=cf5:brand, timestamp=1653667708084, value=dauscher
2019-10-01 00:58:11 UTC	column=cf5:brand, timestamp=1653667708084, value=dauscher
2019-10-01 01:14:09 UTC	column=cf5:brand, timestamp=1653667708084, value=dxracer
2019-10-01 01:20:39 UTC	column=cf5:brand, timestamp=1653667708084, value=dior
2019-10-01 02:17:12 UTC	column=cf5:brand, timestamp=1653667708084, value=dauscher
2019-10-01 02:18:49 UTC	column=cf5:brand, timestamp=1653667708084, value=dewalt
.....	.....
.....	.....
.....	.....

34) hbase(main):003:0> scan 'ecommerce2', {COLUMNS => 'cf5', FILTER => "ValueFilter(=, 'regexstring:el\*be')"}  
 'regexstring:el\*be')"

- Returns key values of the column family 'cf5:brand' in which the name either contain the letters 'el' or 'be'.

ROW	COLUMN+CELL
2019-10-01 02:34:08 UTC	column=cf5:brand, timestamp=1653667708084, value=febest
2019-10-01 03:46:59 UTC	column=cf5:brand, timestamp=1653667708084, value=febest
2019-10-01 03:52:58 UTC	column=cf5:brand, timestamp=1653667708084, value=febest
.....	.....
.....	.....
2019-10-01 04:32:30 UTC	column=cf5:brand, timestamp=1653667708084, value=wellberg
2019-10-01 05:14:59 UTC	column=cf5:brand, timestamp=1653667708084, value=rebellion
.....	.....
.....	.....
2019-10-01 07:35:26 UTC	column=cf5:brand, timestamp=1653667708084, value=mebelev
2019-10-01 08:12:42 UTC	column=cf5:brand, timestamp=1653667708084, value=febest
2019-10-01 09:18:23 UTC	column=cf5:brand, timestamp=1653667708084, value=wellberg
2019-10-01 09:31:20 UTC	column=cf5:brand, timestamp=1653667708084, value=wellberg
2019-10-01 09:32:13 UTC	column=cf5:brand, timestamp=1653667708084, value=wellberg
2019-10-01 10:26:20 UTC	column=cf5:brand, timestamp=1653667708084, value=febest
2019-10-01 10:55:00 UTC	column=cf5:brand, timestamp=1653667708084, value=febest

2019-10-01 10:57:23 UTC                      column=cf5:brand, timestamp=1653667708084, value=ostamebel

.....

.....

.....

35) hbase(main):001:0> scan 'ecommerce2', {COLUMNS => 'cf4', FILTER => "ValueFilter(=, 'binaryprefix:app')"} }

- Returns key values of the column family 'cf4:category\_code' in which the name starts with the string 'app'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf4:category_code, timestamp=1653667708084, value=appliances.environment.water_heater
2019-10-01 00:00:10 UTC	column=cf4:category_code, timestamp=1653667708084, value=apparel.shoes.keds
2019-10-01 00:00:13 UTC	column=cf4:category_code, timestamp=1653667708084, value=appliances.environment.water_heater
2019-10-01 00:00:18 UTC	column=cf4:category_code, timestamp=1653667708084, value=appliances.kitchen.mixer
2019-10-01 00:00:23 UTC	column=cf4:category_code, timestamp=1653667708084, value=appliances.environment.air_heater
2019-10-01 00:00:28 UTC	column=cf4:category_code, timestamp=1653667708084, value=apparel.shoes
2019-10-01 00:00:30 UTC	column=cf4:category_code, timestamp=1653667708084, value=appliances.environment.vacuum
2019-10-01 00:00:31 UTC	column=cf4:category_code, timestamp=1653667708084, value=apparel.shoes.keds
2019-10-01 00:00:33 UTC	column=cf4:category_code, timestamp=1653667708084, value=apparel.shoes
2019-10-01 00:00:35 UTC	column=cf4:category_code, timestamp=1653667708084, value=appliances.kitchen.washer
2019-10-01 00:01:02 UTC	column=cf4:category_code, timestamp=1653667708084, value=appliances.kitchen.mixer
2019-10-01 00:01:21 UTC	column=cf4:category_code, timestamp=1653667708084, value=apparel.shoes.keds
2019-10-01 00:01:23 UTC	column=cf4:category_code, timestamp=1653667708084, value=appliances.kitchen.mixer

.....

.....

.....

36) hbase(main):001:0> scan 'ecommerce2', {STARTROW => '2019-10-01 00:00:19 UTC'}

- Using STARTROW, display only rows from '2019-10-01 00:00:19 UTC' onwards.

ROW	COLUMN+CELL
2019-10-01 00:00:19 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:19 UTC	column=cf2:product_id, timestamp=1653667708084, value=1005135
2019-10-01 00:00:19 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013555631882655	
2019-10-01 00:00:19 UTC	column=cf4:category_code, timestamp=1653667708084,
value=electronics.smartphone	
2019-10-01 00:00:19 UTC	column=cf5:brand, timestamp=1653667708084, value=apple
2019-10-01 00:00:19 UTC	column=cf6:price, timestamp=1653667708084, value=1747.79
2019-10-01 00:00:19 UTC	column=cf7:user_id, timestamp=1653667708084, value=535871217
2019-10-01 00:00:19 UTC	column=cf8:user_session, timestamp=1653667708084, value=c6bd7419-
2748-4c56-95b4-8cec9ff8b80d	
2019-10-01 00:00:20 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:20 UTC	column=cf2:product_id, timestamp=1653667708084, value=4803399
2019-10-01 00:00:20 UTC	column=cf3:category_id, timestamp=1653667708084,
value=2053013554658804075	
2019-10-01 00:00:20 UTC	column=cf4:category_code, timestamp=1653667708084,
value=electronics.audio.headphone	
2019-10-01 00:00:20 UTC	column=cf5:brand, timestamp=1653667708084, value=jbl
2019-10-01 00:00:20 UTC	column=cf6:price, timestamp=1653667708084, value=33.21
2019-10-01 00:00:20 UTC	column=cf7:user_id, timestamp=1653667708084, value=555428858
2019-10-01 00:00:20 UTC	column=cf8:user_session, timestamp=1653667708084, value=8a6afed4-
77f8-40c9-8e76-e062b28216ce	
.....	
.....	
.....	

```
37) hbase(main):002:0> scan 'ecommerce2', {COLUMNS => ['cf2:product_id',
'cf4:category_code'], LIMIT => 3, STARTROW => '2019-10-01 03:10:28 UTC'}
```

- Returns key-value pairs of column families 'cf2' and 'cf4' for three rows starting from row key '2019-10-01 03:10:28 UTC'.

ROW	COLUMN+CELL
2019-10-01 03:10:28 UTC	column=cf2:product_id, timestamp=1653667708084, value=6200702
2019-10-01 03:10:28 UTC	column=cf4:category_code, timestamp=1653667708084,
value=appliances.environment.air_heater	
2019-10-01 03:10:29 UTC	column=cf2:product_id, timestamp=1653667708084, value=1004240
2019-10-01 03:10:29 UTC	column=cf4:category_code, timestamp=1653667708084,
value=electronics.smartphone	
2019-10-01 03:10:30 UTC	column=cf2:product_id, timestamp=1653667708084, value=1004856
2019-10-01 03:10:30 UTC	column=cf4:category_code, timestamp=1653667708084,
value=electronics.smartphone	
3 row(s) in 0.3530 seconds	

```
38) hbase(main):003:0> scan 'ecommerce2', {COLUMNS => ['cf3:category_id', 'cf5:brand'],
STARTROW => '2019-10-01 00:00:20 UTC', STOPROW => '2019-10-01 00:00:26 UTC'}
```

- Returns key-value pairs of column families 'cf3' and 'cf5' starting from row key '2019-10-01 00:00:20 UTC' and stop before '2019-10-01 00:00:26 UTC' (exclusive).

ROW	COLUMN+CELL
2019-10-01 00:00:20 UTC value=2053013554658804075	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:20 UTC	column=cf5:brand, timestamp=1653667708084, value=jbl
2019-10-01 00:00:22 UTC value=2053013561092866779	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:22 UTC	column=cf5:brand, timestamp=1653667708084, value=pulser
2019-10-01 00:00:23 UTC value=2053013552293216471	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:23 UTC	column=cf5:brand, timestamp=1653667708084, value=midea
2019-10-01 00:00:24 UTC value=2061717937420501730	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:24 UTC	column=cf5:brand, timestamp=1653667708084, value=
2019-10-01 00:00:25 UTC value=2053013557225718275	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:25 UTC	column=cf5:brand, timestamp=1653667708084, value=gran-stone

5 row(s) in 0.1370 seconds

```
39) hbase(main):006:0> delete 'ecommerce2', '2019-10-01 00:00:13 UTC', 'cf2:product_id',
1653667708084
hbase(main):007:0> scan 'ecommerce2', {FILTER => "PageFilter(8)"}

```

- Deletes the column family 'cf2:product\_id' present against the row key '2019-10-01 00:00:13 UTC' in the timestamp 1653667708084, and display the first eight rows to confirm it.

0 row(s) in 0.0800 seconds

ROW	COLUMN+CELL
2019-10-01 00:00:13 UTC	column=cf1:event_type, timestamp=1653667708084, value=view
2019-10-01 00:00:13 UTC value=2053013552326770905	column=cf3:category_id, timestamp=1653667708084,
2019-10-01 00:00:13 UTC	column=cf4:category_code, timestamp=1653667708084,
value=appliances.environment.water_heater	
2019-10-01 00:00:13 UTC	column=cf5:brand, timestamp=1653667708084, value=haier

2019-10-01 00:00:13 UTC	column=cf6:price, timestamp=1653667708084, value=102.38
2019-10-01 00:00:13 UTC	column=cf7:user_id, timestamp=1653667708084, value=555444559
2019-10-01 00:00:13 UTC	column=cf8:user_session, timestamp=1653667708084, value=98b88fao-
d8fa-4b9d-8a71-3dd403afab85	
8 row(s) in 0.2050 seconds	

```
40) hbase(main):008:0> deleteall 'ecommerce2', '2019-10-01 00:00:16 UTC'
hbase(main):001:0> scan 'ecommerce2', {FILTER => "RowFilter(=, 'binary:2019-10-01
00:00:16 UTC')"}

```

- *Deletes all the column families (cf1 to cf8) against the row key '2019-10-01 00:00:16 UTC', and display the specified row key to confirm.*

0 row(s) in 0.0210 seconds

ROW

COLUMN+CELL

0 row(s) in 29.7740 seconds

```
41) hbase(main):041:0> scan 'ecommerce2', {COLUMNS => ['cf7','cf8'], LIMIT => 6}

```

- *Returns the column families 'cf7' and 'cf8' present against the starting 6 row keys.*

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf7:user_id, timestamp=1653667708084, value=554748717
2019-10-01 00:00:00 UTC	column=cf8:user_session, timestamp=1653667708084, value=9333dfbd-
b87a-4708-9857-6336556b0fcc	
2019-10-01 00:00:01 UTC	column=cf7:user_id, timestamp=1653667708084, value=550050854
2019-10-01 00:00:01 UTC	column=cf8:user_session, timestamp=1653667708084, value=7c90fc70-
0e80-4590-96f3-13c02c18c713	
2019-10-01 00:00:04 UTC	column=cf7:user_id, timestamp=1653667708084, value=535871217
2019-10-01 00:00:04 UTC	column=cf8:user_session, timestamp=1653667708084, value=c6bd7419-
2748-4c56-95b4-8cec9ff8b80d	
2019-10-01 00:00:05 UTC	column=cf7:user_id, timestamp=1653667708084, value=512742880
2019-10-01 00:00:05 UTC	column=cf8:user_session, timestamp=1653667708084, value=0d0d91c2-
c9c2-4e81-90a5-86594dec0db9	
2019-10-01 00:00:08 UTC	column=cf7:user_id, timestamp=1653667708084, value=550978835
2019-10-01 00:00:08 UTC	column=cf8:user_session, timestamp=1653667708084, value=6280d577-
25c8-4147-99a7-abc6048498d6	
2019-10-01 00:00:10 UTC	column=cf7:user_id, timestamp=1653667708084, value=520571932
2019-10-01 00:00:10 UTC	column=cf8:user_session, timestamp=1653667708084, value=ac1cd4e5-
a3ce-4224-a2d7-ff660a105880	
6 row(s) in 0.0590 seconds	

42) hbase(main):043:0> scan 'ecommerce2', {ROWPREFIXFILTER => '2', FILTER => "QualifierFilter(=, 'binary:price')"} }

- Returns the key values of column 'price' present against each row key which has a prefix of '2'.

ROW	COLUMN+CELL
2019-10-01 00:00:00 UTC	column=cf6:price, timestamp=1653667708084, value=33.20
2019-10-01 00:00:01 UTC	column=cf6:price, timestamp=1653667708084, value=251.74
2019-10-01 00:00:04 UTC	column=cf6:price, timestamp=1653667708084, value=1081.98
2019-10-01 00:00:05 UTC	column=cf6:price, timestamp=1653667708084, value=908.62
2019-10-01 00:00:08 UTC	column=cf6:price, timestamp=1653667708084, value=41.16
2019-10-01 00:00:10 UTC	column=cf6:price, timestamp=1653667708084, value=102.71
2019-10-01 00:00:11 UTC	column=cf6:price, timestamp=1653667708084, value=900.64
2019-10-01 00:00:13 UTC	column=cf6:price, timestamp=1653667708084, value=102.38
2019-10-01 00:00:15 UTC	column=cf6:price, timestamp=1653667708084, value=35.79
2019-10-01 00:00:17 UTC	column=cf6:price, timestamp=1653667708084, value=357.79
2019-10-01 00:00:18 UTC	column=cf6:price, timestamp=1653667708084, value=58.95
2019-10-01 00:00:19 UTC	column=cf6:price, timestamp=1653667708084, value=1747.79
2019-10-01 00:00:20 UTC	column=cf6:price, timestamp=1653667708084, value=33.21
2019-10-01 00:00:22 UTC	column=cf6:price, timestamp=1653667708084, value=921.49
2019-10-01 00:00:23 UTC	column=cf6:price, timestamp=1653667708084, value=47.62
2019-10-01 00:00:24 UTC	column=cf6:price, timestamp=1653667708084, value=151.87
.....	
.....	
.....	

## Snapshot Shell Commands:

1) hbase(main):004:0> snapshot 'ecommerce2', 'ecomm\_data'

- *Creates a snapshot named 'ecomm\_data' of the specified table 'ecommerce2'.*

0 row(s) in 2.2580 seconds

2) hbase(main):007:0> list\_snapshots

- *List all the snapshots taken. The only snapshot 'ecomm\_data' will appear.*

SNAPSHOT	TABLE + CREATION TIME
ecomm_data	ecommerce2 (Sun May 29 13:22:13 +0000 2022)

1 row(s) in 0.0420 seconds

=> ["ecomm\_data"]

3) hbase(main):005:0> clone\_snapshot 'ecomm\_data', 'ecommerce\_data'  
hbase(main):006:0> list

- *Creates a new table 'ecommerce\_data' and clone the snapshot content into it.*

0 row(s) in 8.4430 seconds

TABLE
ecommerce2
ecommerce_data

2 row(s) in 0.0180 seconds

4) hbase(main):008:0> delete\_snapshot 'ecomm\_data'  
hbase(main):009:0> list\_snapshots

- *Deletes the snapshot 'ecomm\_data' which was created earlier.*

0 row(s) in 0.0980 seconds

SNAPSHOT	TABLE + CREATION TIME
----------	-----------------------

0 row(s) in 0.0110 seconds

=> []

## Replication Commands:

```
1) hbase(main):002:0> add_peer '1', "server1.cie.com:2181:/hbase"
hbase(main):001:0> add_peer '2', "zk1,zk2:2182:/hbase-prod"
```

- *Adding peers '1' and '2' to cluster to replicate the data.*

0 row(s) in 0.2920 seconds

0 row(s) in 0.3860 seconds

```
2) hbase(main):002:0> list_peers
```

- *List all replication peer clusters.*

```
PEER_ID CLUSTER_KEY STATE TABLE_CFS
```

```
1 server1.cie.com:2181:/hbase ENABLED
```

```
2 zk1,zk2:2182:/hbase-prod ENABLED
```

2 row(s) in 0.0510 seconds

```
3) hbase(main):014:0> list_replicated_tables
```

- *List all the tables and column families replicated from the cluster.*

TABLE:COLUMNFAMILY	ReplicationType
ecommerce2:cf1	GLOBAL
ecommerce2:cf2	GLOBAL
ecommerce2:cf3	GLOBAL
ecommerce2:cf4	GLOBAL
ecommerce2:cf5	GLOBAL
ecommerce2:cf6	GLOBAL
ecommerce2:cf7	GLOBAL
ecommerce2:cf8	GLOBAL
ecommerce2:event_time	GLOBAL

9 row(s) in 0.0330 seconds



4) hbase(main):027:0> disable\_peer '1'  
hbase(main):028:0> list\_peers

- *Stops the replication stream to the specified peer cluster '1', and returns the list to confirm.*

0 row(s) in 0.0580 seconds

PEER\_ID CLUSTER\_KEY STATE TABLE\_CFS

1 server1.cie.com:2181:/hbase DISABLED

2 zk1,zk2:2182:/hbase-prod ENABLED

2 row(s) in 0.0300 seconds

5) hbase(main):029:0> enable\_peer '1'  
hbase(main):030:0> list\_peers

- *Restarts the replication stream to the specified peer cluster '1', and returns the list to confirm.*

0 row(s) in 0.0910 seconds

PEER\_ID CLUSTER\_KEY STATE TABLE\_CFS

1 server1.cie.com:2181:/hbase ENABLED

2 zk1,zk2:2182:/hbase-prod ENABLED

2 row(s) in 0.0160 seconds

6) hbase(main):031:0> remove\_peer '2'  
hbase(main):032:0> list\_peers

- *Stops the specified replication stream '2' and deletes all the meta information kept about it. Returns the list to confirm.*

0 row(s) in 0.0720 seconds

PEER\_ID CLUSTER\_KEY STATE TABLE\_CFS

1 server1.cie.com:2181:/hbase ENABLED

1 row(s) in 0.0090 seconds